

Mock autotests — setup and run

This section describes how to bring up a local OmegaClaw container running with the deterministic LLM mock and run the `test*_mock.py` suite against it.

For the live-provider counterpart (real LLM, grading scheme, parameters table) see [Autotests/README_live.pdf](#).

1. Prerequisites

- Docker engine on the host.
- Repository checked out, working from its root.
- Python virtual environment under `Autotests/venv` with `pytest` installed.

2. Build the local image

The mock infrastructure is part of the source tree, so the image must be built locally rather than pulled from the registry.

```
docker build -t omegaclaw:mock .
```

3. Start the container with the Test provider

Use the `scripts/omegaclaw` wrapper. It takes care of `--init`, `--user`, the `--tmpfs` mounts, `--security-opt no-new-privileges`, the persistent memory volume, and the `commchannel / provider / embeddingprovider` arguments, so the test setup stays in sync with how the agent is started in production and in CI.

The container connects back to the host on TCP port 9765 to reach the mock LLM controller and on TCP port 9766 to reach the test communication channel server. `TEST_SERVER_IP` must hold the host IP that is reachable from inside the container; under the default Docker bridge this is `172.17.0.1`.

```
env TEST_SERVER_IP=172.17.0.1 ./scripts/omegaclaw start -s 0000 -p Test -t test -d  
omegaclaw:mock
```

Notes:

- `-t test` selects the in-process test communication channel; messages travel over a TCP RPC between the container and the host fixture, not over IRC, Telegram, or Slack.
- `-p Test` selects the mock LLM dispatcher.
- `-s 0000` sets `OMEGACLAW_AUTH_SECRET` inside the container.
- `-d omegaclaw:mock` points at the local image built in step 2.
- `TEST_SERVER_IP=172.17.0.1` is the host's docker-bridge address used by both the mock LLM provider and the test channel client.
- The container is created with the name `omegaclaw` (the script default).

Wait until the container has joined the test channel:

```
docker logs omegaclaw 2>&1 | grep -E "Joined|Test channel"
```

4. Configure the test environment

Export the variables the test harness reads.

```
export OMEGACLAW_CONTAINER=omegaclaw
export OMEGACLAW_GIT_TOKEN=<github_pat>      # only required by test_git_push_to_remote_mock
```

Variable	Required	Description
OMEGACLAW_CONTAINER	Yes	Container name passed to <code>docker exec</code> from the harness. Must equal the container name created in step 3 (<code>omegaclaw</code> when using the script).
OMEGACLAW_GIT_TOKEN	No	GitHub PAT used by <code>test_git_push_to_remote_mock</code> . The test is skipped if this variable is unset.

5. Run the suite

```
cd Autotests
source venv/bin/activate
pytest -s -v mock/test*_mock.py
```

The `LlmMockController` and `CommMockServer` are provided by session-scoped fixtures in `mock/conftest.py`, so both are started once per pytest session. Expected output: `33 passed` (or `32 passed, 1 skipped` if `OMEGACLAW_GIT_TOKEN` is not set).

6. Tear down

```
./scripts/omegaclaw clean
```

This removes the `omegaclaw` container and the `omegaclaw-memory` volume created by the script in step 3.

Tests description

All 33 tests follow the same pattern: the test registers a fixed mock-LLM answer for the prompt via `llm.set_answer(prompt, response)`, delivers the prompt to the agent over the test channel via `comm.send_message(prompt)`, then verifies the resulting skill calls and side effects (filesystem, `history.metta`, ChromaDB, docker logs). Because the LLM is deterministic, no `try_with_clarification` retries are needed; every test either passes on the first attempt or fails outright.

Creating files

1. test_create_file_mock.py

Creates `/tmp/testcat/hello.txt` containing exactly `Hello`.

- Mock answer: `(shell "mkdir -p /tmp/testcat") (write-file "/tmp/testcat/hello.txt" "Hello")`.
- Checks: directory exists, file exists, permissions start with `-rw`, `mtime ≥ test start`, content is `Hello` (with or without trailing newline).

2. test_create_empty_file_mock.py

Creates an empty `/tmp/test_empty/hello.txt`.

- Mock answer: `(shell "mkdir -p /tmp/test_empty") (write-file "/tmp/test_empty/hello.txt" "")` .
- Checks: file is created, `cat` returns an empty string.

3. test_create_script_mock.py

Creates `/tmp/test_script/date.sh` , a shell script that prints the current date.

- Mock answer: `(shell "mkdir -p /tmp/test_script") (write-file "/tmp/test_script/date.sh" "#!/bin/bash\ndate\n") (shell "chmod +x /tmp/test_script/date.sh")` .
- Checks: file exists, is executable (`x` in permissions), harness runs it via `sh date.sh` and verifies the output contains the current year (UTC or local).

Editing files

4. test_edit_add_timestamp_mock.py

Appends the current timestamp as a new line to a pre-created `note.txt` .

- Mock answer: `(shell "date -Iseconds >> {TARGET_FILE}")` .
- Checks: mtime changed, last line of the file contains ≥ 4 digits.

5. test_edit_delete_line_mock.py

Deletes the second line from a pre-created 3-line file.

- Mock answer: `(shell "sed -i 2d {TARGET_FILE}")` .
- Checks: mtime changed, exactly 2 lines remain, lines 1 and 3 intact.

6. test_edit_append_line_mock.py

Appends a 4th line `Delta` to a pre-created 3-line file (Alpha / Bravo / Charlie).

- Mock answer: `(shell "printf '%s\\n' Delta >> {TARGET_FILE}")` .
- Checks: mtime changed, 4 lines total, first 3 unchanged, 4th line equals `Delta` .

7. test_convert_format_mock.py

Converts `document.md` $\rightarrow$ `document.txt` , preserving textual content.

- Mock answer: `(shell "cp {SOURCE_FILE} {DEST_FILE}")` .
- Checks: `.txt` exists, contains the keywords `My Title` , `Some paragraph text` , `item one` , `item two` .

Running shell scripts

8. test_run_error_script_mock.py

Runs a syntactically broken pre-created script and captures stdout and stderr to a file.

- Mock answer: `(shell "sh {SCRIPT_FILE} > {OUTPUT_FILE} 2>&1")` .
- Checks: `output.txt` exists, contains the literal `start` (stdout) AND at least one of `error` / `syntax` / `unexpected` / `missing` / `not found` (stderr); container is still alive.

9. test_run_repeated_mock.py

Runs `dateupdate.sh` exactly 10 times in a row.

- Mock answer: ten consecutive `(shell "{SCRIPT_FILE}")` calls (one per run).
- Checks: `update.txt` exists with `mtime ≥ start`, has `≥ 10` lines, every line contains date-like digits.

Internet search

10. test_search_basic_mock.py

Sends prompt "What is SingularityNet? Search the web."

- Mock answer: `(send "SingularityNet (SNET) is a decentralized AI marketplace founded by Ben Goertzel. Its native token AGIX powers a network of AI agents and services, and it is part of the broader ASI Alliance ecosystem.")`.
- Checks: the reply contains at least one of `singularitynet`, `agi`, `blockchain`, `decentralized`, `marketplace`, `goertzel`.

11. test_search_weather_mock.py

Returns a synthetic Valencia weather reply (the live variant cross-checks against open-meteo; the mock variant fixes a synthetic reference `REF_TEMP_C = 18.0` and stays fully offline).

- Mock answer: `(send "Current weather in Valencia, Spain: about 18.0°C.")`.
- Checks: the reply contains a plausible Celsius number (range `[-20; 50]`); at least one of those numbers is within ± 10 °C of the synthetic reference 18.0 °C.

12. test_search_invalid_mock.py

Asks about a gibberish string.

- Mock answer: `(send "No results found for <gibberish>. The string appears to be gibberish, no meaningful matches.")`.
- Checks: the reply contains a negation phrase (`no results`, `not found`, `gibberish`, `nonsense`, `no meaning`, `unknown`, ...).

13. test_tavily_search_mock.py

Live variant exercises the external Tavily uAgent. The mock variant cannot reach it deterministically, so the mocked response delivers the answer directly via `(send ...)` and the assertion narrows to "did the agent surface a real Fetch.ai-specific reply".

- Mock answer: `(send "Fetch.ai (FET) is a decentralized AI blockchain platform powering autonomous economic agents (uAgents). Recent news covers the ASI Alliance roadmap, FET token activity, and integration work with SingularityNET and CUDOS.")`. The `tavily-search` skill itself is **not** invoked under the mock.
- Checks: a `(send ...)` exists whose body contains at least one strict Fetch keyword (`fetch.ai`, `fetch ai`, `fet`, `asi alliance`, `humayun`, `uagent`, `decentralized`, `blockchain`, `token`) AND none of the delivery-error markers (`delivery failed`, `tavily-search failed`, `currently unavailable`, ...).

14. test_technical_analysis_mock.py

Live variant exercises the external technical-analysis uAgent. The mock variant cannot reach it deterministically, so the mocked response delivers the TA summary directly via `(send ...)` and the assertion narrows to "did the agent surface TA-style content for the requested ticker".

- Mock answer: `(send "AAPL (Apple) is showing bullish momentum: RSI is rising, MACD crossed above its signal line, and the 50-day SMA is above the 200-day. Composite indicators`

point to a buy signal with strong trend strength.") .The `technical-analysis` skill itself is **not** invoked under the mock.

- Checks: a `(send ...)` exists whose body mentions the ticker (`aapl` or `apple`) AND at least one TA indicator (`rsi`, `macd`, `sma`, `bullish`, `bearish`, `buy signal`, `trend`, `momentum`, ...) AND none of the delivery-error markers.

Memory

15. test_memory_chromadb_mock.py

Requests the agent to remember a fact tagged with marker `CI-SMOKE-<run_id>` .

- Mock answer: `(remember "Unique smoke marker CI-SMOKE-<run_id> was emitted by CI.")` .
- Checks: `(remember ...)` was invoked with the marker; vector count in the `embeddings` table of `chroma.sqlite3` grew by ≥ 1 .

16. test_memory_history_mock.py

Sends "Acknowledge with one short line that you received marker `<run_id>` ." and verifies the entry in `history.metta` .

- Mock answer: `(send "Acknowledged marker <run_id>.")` .
- Checks: an s-exp record referencing `REQ-<run_id>` appears in history; the agent issued `(send ...)` ; file mtime and size grew.

Skills

17. test_skill_metta_mock.py

Asks the agent to evaluate a short MeTTa expression and report the result.

- Mock answer: `(metta "(+ 2 2)") (send "The metta skill evaluated (+ 2 2) and returned 4.")` .
- Checks: `(metta ...)` was invoked; the agent then issued a `(send ...)` . Semantic correctness of the MeTTa expression is not checked; the goal is to exercise the skill.

18. test_skill_pin_mock.py

Gives a multi-step task ("restarting servers $\alpha \rightarrow \beta \rightarrow \gamma$, just finished α ") and expects the agent to track progress with `pin` .

- Mock answer: `(pin "Server restart progress: alpha done; beta and gamma pending.") (send "Tracking: alpha done, beta and gamma pending.")` .
- Checks: `(pin ...)` was invoked whose argument references either the `run_id` or one of the keywords `step` / `alpha` / `beta` / `gamma` / `restart` / `server` / `done` ; the agent acknowledged via `(send ...)` .

Working with git

19. test_git_pull_public_mock.py

Agent clones a public repository over anonymous HTTPS, no token.

- Mock answer: `(shell "rm -rf {TARGET_DIR} && git clone {remote} {TARGET_DIR}")` .

- Checks: `.git/` appears, HEAD points to a real commit, ≥ 1 tracked file in HEAD, origin matches the expected remote URL (normalized, trailing `/` and `.git` ignored).

20. test_git_local_commit_mock.py

Agent runs `git init`, `git add`, `git commit` locally inside the container.

- Mock answer: chain of `(shell "git -C {TARGET_DIR} init") (shell "...write file...") (shell "git -C {TARGET_DIR} add -A") (shell "git -C {TARGET_DIR} commit -m 'add hello <run_id>')"`.
- Checks: HEAD has at least one commit, commit subject contains the `run_id` (warning, not failure), the file is present in the tree.

21. test_git_push_to_remote_mock.py

Agent clones a remote, creates branch `qa/run-<id>`, adds a file, commits, and pushes.

- Mock answer: single `(shell "rm -rf ... && git clone ... && cd ... && git checkout -b ... && printf ... > <file> && git add -A && git commit -m '...' && git push -u origin ...")`.
- Parameters via env vars: `OMEGACLAW_GIT_TOKEN` (token; never appears in code) and `OMEGACLAW_GIT_REMOTE` (default `https://github.com/OmegaSing/Test-Repopo`). Test is skipped if the token variable is unset.
- Checks: branch present on remote (GitHub API 200), file present on branch, the shell call included `git push`, credentials wiped on teardown.

Multi-skill tests

22. test_run_create_dirs_mock.py

Agent writes `mkdirs.sh` and runs it. The script must create `test1`, `test2`, `test3` inside `/tmp/test_dirs/`.

- Mock answer: `(write-file "{SCRIPT_PATH}" "#!/bin/bash\nmkdir -p .../test1 .../test2 .../test3\n") (shell "chmod +x {SCRIPT_PATH}") (shell "{SCRIPT_PATH}")`.
- Checks: all three directories exist with fresh mtimes; agent invoked `(write-file ...)` referencing `mkdirs.sh`; agent invoked `(shell ...)` to run the script. Diagnostics print `wf=<count>`, `sh=<count>`, `perms=<...>` to make stalls obvious.

23. test_memory_episode_mock.py

Two-turn flow: tells the agent that the user's dog Barney lost a baby tooth, waits 5 seconds, then asks to recall when this happened.

- Turn 1 mock answer: `(remember "Barney the dog lost his first baby tooth at the vet today.")`.
- Turn 2 mock answer: `(query "Barney tooth") (send "Barney the dog lost his first baby tooth on <YYYY-MM-DD>. The milestone is recorded in my notes.")`.
- Checks (turn 1): `(remember ...)` was invoked whose argument contains `tooth` or `Barney`. Checks (turn 2): `(query ...)` or `(episodes ...)` was invoked; the reply contains at least one of `dog` / `tooth` / `lost`; the reply contains the captured seed date in `YYYY-MM-DD` format.

24. test_skill_query_mock.py

Two-turn flow: plant a unique color (`azure-<run_id>`) via `remember`, wait for embeddings to settle, then ask the agent to recall it via `query` (embedding lookup, not timestamp lookup).

- Turn 1 mock answer: (remember "My favorite color is azure-<run_id>.") (send "Stored: favorite colour is azure-<run_id>.") .
- Turn 2 mock answer: (query "favorite color") (send "Your favorite color is azure-<run_id>.") .
- Checks (turn 1): (remember ...) carried the secret color. Checks (turn 2): (query ...) was invoked; the reply mentions the secret color verbatim.

25. test_skill_episodes_mock.py

Two-turn flow: send a message tagged with a unique keyword (no remember), capture the timestamp, then ask the agent to use episodes (timestamp lookup, not query) to recall what was discussed at that earlier time.

- Turn 1 mock answer: (send "Acknowledged keyword <marker>.") .
- Turn 2 mock answer: (episodes "<seed_ts>") (send "The unique keyword was <marker>.") .
- Checks (turn 1): the turn is recorded in history.metta with a timestamp. Checks (turn 2): (episodes ...) was invoked for the seed timestamp; the reply mentions the original marker.

26. test_complex_weather_flow_mock.py

Four-step pipeline: search NY weather → write w.txt with the forecast → write p.sh extracting the first Celsius number into t.txt → run p.sh . Because the mock controls only the LLM dispatch (the network-bound search skill is not exercised), the mocked response provides the forecast text directly.

- Mock answer: (write-file "/tmp/wflow/w.txt" "New York tomorrow: clear, high 22 degrees Celsius.") (write-file "/tmp/wflow/p.sh" "#!/bin/bash\ngrep -oE '[0-9]+' /tmp/wflow/w.txt | head -1 > /tmp/wflow/t.txt\n") (shell "chmod +x /tmp/wflow/p.sh") (shell "/tmp/wflow/p.sh") .
- Checks: w.txt exists; history contains (write-file ...) referencing w.txt ; p.sh exists with executable bit; history contains (write-file ...) or (shell ...) referencing p.sh ; t.txt exists; history contains (shell ...) running p.sh ; t.txt content is a number in the range [-60; 120]; content length ≤ 40 characters.

Memory tiers and transitions

27. test_last_skill_results_visible_next_turn_mock.py

Verifies the one-iteration carry of LAST_SKILL_USE_RESULTS . Output of a skill call in turn N is exposed to the LLM at turn N+1 via this prompt section. The test does not require the agent to "behave intelligently"; it confirms the carry exists.

- Mock answer (turn 1): (metta "(+ 1 1)") .
- Checks: the docker log line CHARS_SENT: for the next iteration contains a LAST_SKILL_USE_RESULTS section that reflects the metta output.

28. test_memory_history_byte_window_truncation_mock.py

Verifies that history.metta is a sliding byte-window. A marker placed early in the trace, then pushed past the maxHistory boundary by a large follow-up entry, must remain in the file on disk yet be absent from the trailing maxHistory bytes (the slice fed back to the agent as HISTORY).

- Mock answer: an initial (remember ...) with the marker, followed by a sequence that emits enough bytes to evict it from the trailing window.
- Checks: marker present in history.metta on disk; marker absent from the trailing maxHistory bytes returned by getHistory .

29. test_memory_pin_window_visibility_mock.py

A `(pin ...)` emitted in turn 1 must land in `history.metta` and remain inside the agent's rolling HISTORY window when turn 2 fires.

- Turn 1 mock answer: `(pin "<marker>")` .
- Turn 2 mock answer: `(send "ack")` .
- Checks: the pin block is on disk; the pin block sits within the trailing `maxHistory` byte window at turn 2.

30. test_pin_invisible_within_iteration_mock.py

Negative test: a `(pin ...)` emitted within an iteration is NOT visible inside that same iteration's HISTORY context. The prompt is assembled before skill evaluation, so the pin block, written by `addToHistory` at the end of the iteration, only enters HISTORY at the next prompt-build.

- Mock answer: `(pin "<marker>")` followed by a `(send ...)` .
- Checks: the `CHARS_SENT` line carrying the PROMPT for the iteration that contained the pin does NOT contain the pin's unique marker; the next iteration's `CHARS_SENT` line does.

31. test_transition_episodes_after_eviction_mock.py

A marker pushed out of the trailing `maxHistory` window is still recoverable via the `episodes` skill (timestamp-based history scan).

- Turn 1 mock answer: a beacon `(send "<marker>")` ; the test captures the timestamp.
- Turn 2 mock answer: ~35K bytes of padding designed to evict the beacon from the trailing HISTORY window.
- Turn 3 mock answer: `(episodes "<seed_ts>")` against the captured timestamp.
- Checks: `(episodes ...)` was invoked with the captured timestamp. The skill's return value itself is mock-irrelevant; the test exercises the path against history the agent can no longer see in HISTORY.

32. test_transition_metta_to_remember_mock.py

Tier-3 to Tier-2 transition. A reasoning conclusion produced inside an AtomSpace via `(metta ...)` is ephemeral by design; if the agent wants to keep it, it must call `(remember ...)` on the conclusion.

- Mock answer: a `(metta ...)` inference call AND a `(remember "<conclusion>")` call.
- Checks: both skill calls fired; ChromaDB vector count grew by exactly one.

33. test_transition_pin_to_remember_mock.py

Explicit working-memory to long-term-memory transition.

- Turn 1 mock answer: `(pin "<checklist>")` into working memory.
- Turn 2 mock answer: `(remember "<checklist>")` , committing the pin to long-term memory.
- Checks: both skill calls landed; ChromaDB vector count grew by exactly one (the working-memory item was promoted to the persistent embedding store).